

# CPU capacity forecasting & node failure scenarios

By: Karl Arao

## CPU capacity

Below are the sizing scenarios (workload growth and node fail) and forecasting for headroom expiry days (# of days to reach 75% utilization).

- The current CPU utilization is at 39% clusterwide (49.92 out of 128 CPU threads total capacity)



## Node Failure Scenarios

- First, let's do node failure scenarios using the sizing worksheet [https://github.com/karlarao/sizing\\_worksheet](https://github.com/karlarao/sizing_worksheet)

### Overall Utilization:

Total CPUs & pct% USED	Total Mem GB & pct% USED
49.92 CPUs 39.0%	92 GB 9.0%

### Recommended Hardware:

Equivalent compute nodes
1.6 nodes

Total Host CPUs	CPU Utilization	DB name	PGA Max (G)	SGA Max (G)	node 1	node 2	node 3	node 4
					1 instance 39% cpu used 9% mem used 22G HPages	1 instance 39% cpu used 9% mem used 22G HPages	1 instance 39% cpu used 9% mem used 22G HPages	1 instance 39% cpu used 9% mem used 22G HPages
128	39.00%	PDOIE01	3	20	P	P	P	P

- Replacing the P on the 4th node with F will fail the node. It shows 52% on the remaining three nodes. Here you can afford to lose one more node. And this is also assuming all of the nodes are well balanced in terms of CPU workload.

### Overall Utilization:

Total CPUs & pct% USED	Total Mem GB & pct% USED
49.92 CPUs 39.0%	69 GB 6.7%

### Recommended Hardware:

Equivalent compute nodes
1.6 nodes

Total Host CPUs	CPU Utilization	DB name	PGA Max (G)	SGA Max (G)	node 1	node 2	node 3	node 4
					1 instance 52% cpu used 9% mem used 22G HPages	1 instance 52% cpu used 9% mem used 22G HPages	1 instance 52% cpu used 9% mem used 22G HPages	0 instance 0% cpu used 0% mem used 0G HPages
128	39.00%	PDOIE01	3	20	P	P	P	F

- 2 nodes fail, it shows 78% on the remaining two nodes

### Overall Utilization:

Total CPUs & pct% USED	Total Mem GB & pct% USED
49.92 CPUs 39.0%	46 GB 4.5%

### Recommended Hardware:

Equivalent compute nodes
1.6 nodes

Total Host CPUs	CPU Utilization	DB name	PGA Max (G)	SGA Max (G)	node 1	node 2	node 3	node 4
					1 instance 78% cpu used 9% mem used 22G HPages	1 instance 78% cpu used 9% mem used 22G HPages	0 instance 0% cpu used 0% mem used 0G HPages	0 instance 0% cpu used 0% mem used 0G HPages
128	39.00%	PDOIE01	3	20	P	P	F	F

- With all 4 nodes running change the 39% to 75% to simulate a workload growth. Let's say the current workload continue to increase.

### Overall Utilization:

Total CPUs & pct% USED	Total Mem GB & pct% USED
96 CPUs 75.0%	92 GB 9.0%

### Recommended Hardware:

Equivalent compute nodes
3 nodes

Total Host CPUs	CPU Utilization	DB name	PGA Max (G)	SGA Max (G)	node 1	node 2	node 3	node 4
					1 instance 75% cpu used 9% mem used 22G HPages	1 instance 75% cpu used 9% mem used 22G HPages	1 instance 75% cpu used 9% mem used 22G HPages	1 instance 75% cpu used 9% mem used 22G HPages
128	75.00%	PDOIE01	3	20	P	P	P	P

- Then fail a node. It shows 100% on the remaining three nodes. With 75% clusterwide utilization we can't even afford to lose a node. At 75% it's already a high utilization, and at 100% everything will be very unstable.

### Overall Utilization:

Total CPUs & pct% USED	Total Mem GB & pct% USED
96 CPUs 75.0%	69 GB 6.7%

### Recommended Hardware:

Equivalent compute nodes
3 nodes

Total Host CPUs	CPU Utilization	DB name	PGA Max (G)	SGA Max (G)	node 1	node 2	node 3	node 4
					1 instance 100% cpu used 9% mem used 22G HPages	1 instance 100% cpu used 9% mem used 22G HPages	1 instance 100% cpu used 9% mem used 22G HPages	0 instance 0% cpu used 0% mem used 0G HPages
128	75.00%	PDOIE01	3	20	P	P	P	F

- 2 nodes fail, it shows 150% on the remaining two nodes

### Overall Utilization:

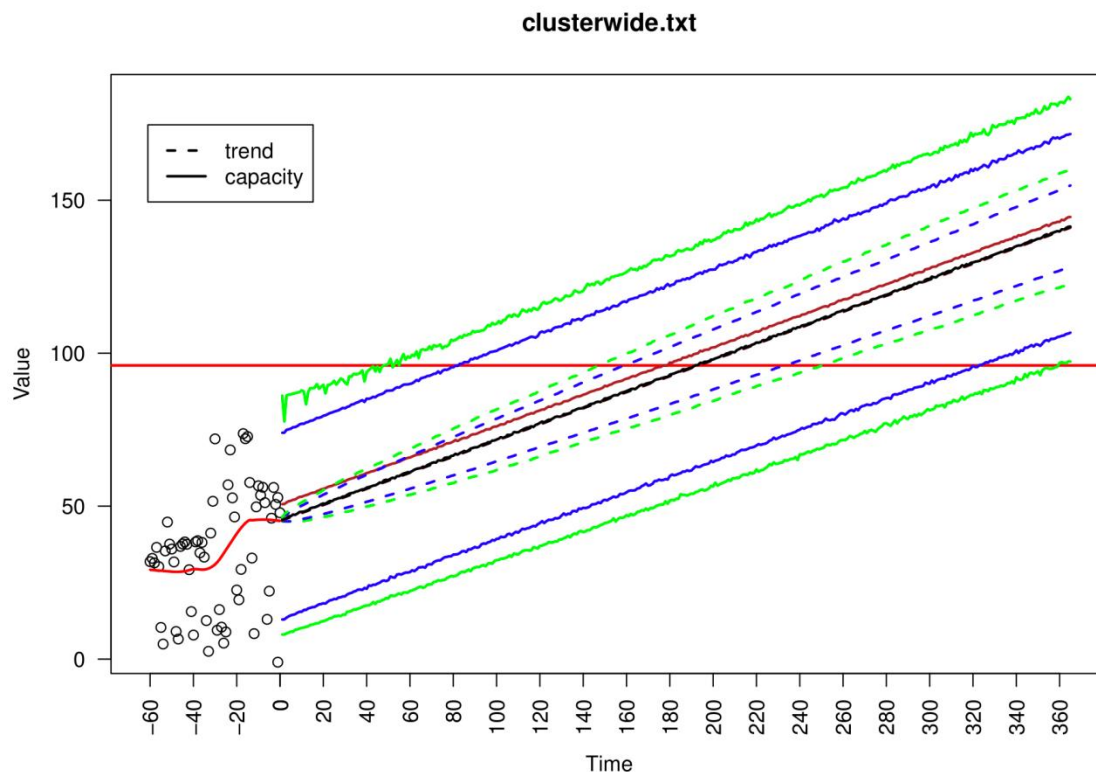
Total CPUs & pct% USED	Total Mem GB & pct% USED
96 CPUs 75.0%	46 GB 4.5%

### Recommended Hardware:

Equivalent compute nodes
3 nodes

Total Host CPUs	CPU Utilization	DB name	PGA Max (G)	SGA Max (G)	node 1	node 2	node 3	node 4
					1 instance 150% cpu used 9% mem used 22G HPages	1 instance 150% cpu used 9% mem used 22G HPages	0 instance 0% cpu used 0% mem used 0G HPages	0 instance 0% cpu used 0% mem used 0G HPages
128	75.00%	PDOIE01	3	20	P	P	F	F

## Forecasting headroom expiry



Based on the clusterwide CPU demand forecast we will reach the 75% of the clusterwide CPU (that's 128CPUs x .75 = 96 CPUs) around the 48th day (about 1.5months) at 99% forecast quantile and 84<sup>th</sup> day (about 2.7 months) at 95% forecast quantile.

To summarize, the current 39% clusterwide CPU utilization is still acceptable given that 1 node failure can be tolerated. But there should be a remediation/tuning activity done immediately so the upward growth will be flat or decrease. If the capacity remediation/tuning is not successful or can't be done then adding one more nodes is a must.

### Commands used:

```
E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R
clusterwide.txt 96 30 .99 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] 192

$quantiles
      quantile headroom
[1,]      0.99      48

E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R
clusterwide.txt 96 30 .95 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] 193

$quantiles
      quantile headroom
[1,]      0.95      84
```

## Appendices


### Setting up the data

The time series CPU data used here came from the top events script of run\_awr tool  
[https://github.com/karlarao/run\\_awr-quickextract/blob/master/run\\_awr\\_topevents.sql](https://github.com/karlarao/run_awr-quickextract/blob/master/run_awr_topevents.sql)

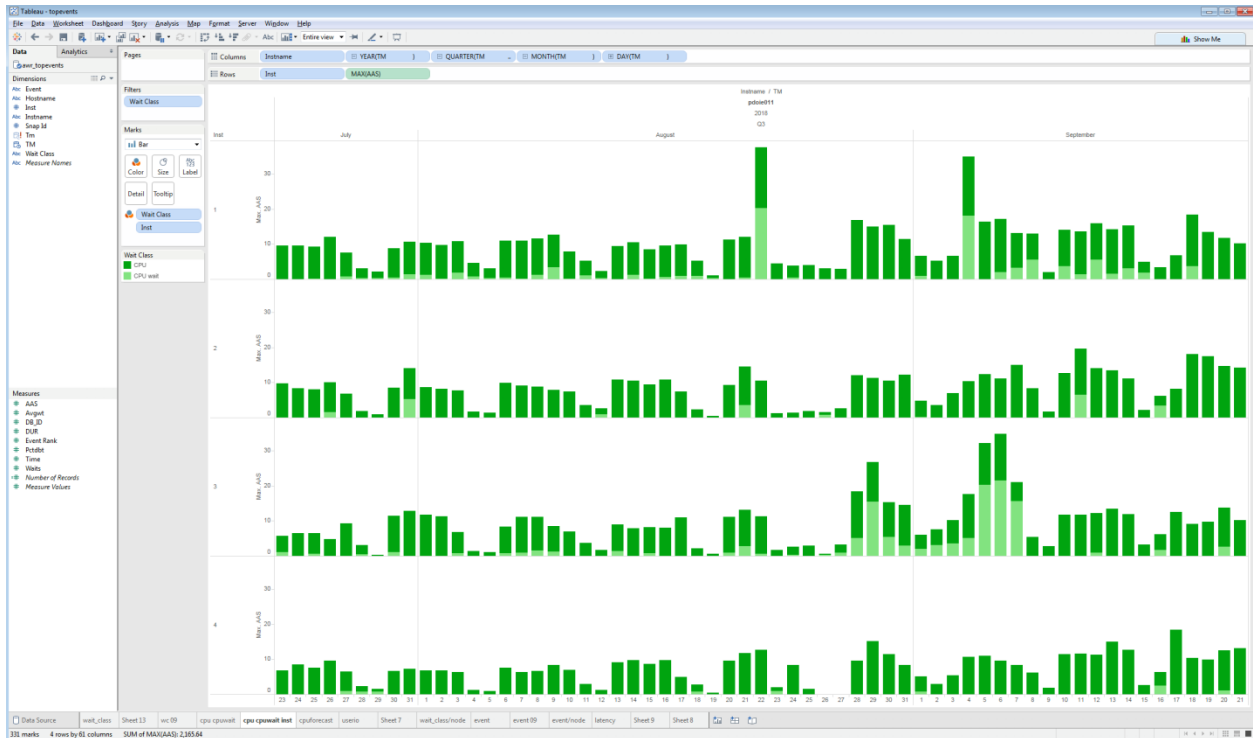
The host OS CPU utilization data can also be used which can be gathered using  
[https://github.com/karlarao/run\\_awr-quickextract/blob/master/run\\_awr\\_cpuwl.sql](https://github.com/karlarao/run_awr-quickextract/blob/master/run_awr_cpuwl.sql)

I use both data whenever I do sizing and forecasting exercise but for this example I'm only showing the top events.

- Visualize the clusterwide CPU utilization



- Split the data by node and filter just the CPU and CPU wait



- Rotate the data and split it by Instance and Wait Class



- [illegible]



- Create a text file for each node and clusterwide

clusterwide.txt	node1.txt	node2.txt	node3.txt	node4.txt
1 Date CPU_AAS	1 Date Value	1 Date Value	1 Date Value	1 Date Value
2 23-Jul 31.89	2 23-Jul 9.60	2 23-Jul 9.75	2 23-Jul 5.74	2 23-Jul 6.81
3 24-Jul 32.92	3 24-Jul 9.65	3 24-Jul 8.29	3 24-Jul 6.45	3 24-Jul 8.52
4 25-Jul 31.48	4 25-Jul 9.36	4 25-Jul 8.01	4 25-Jul 6.55	4 25-Jul 7.57
5 26-Jul 36.53	5 26-Jul 12.09	5 26-Jul 10.05	5 26-Jul 4.85	5 26-Jul 9.54
6 27-Jul 30.25	6 27-Jul 7.54	6 27-Jul 6.85	6 27-Jul 9.31	6 27-Jul 6.55
7 28-Jul 10.34	7 28-Jul 3.14	7 28-Jul 1.85	7 28-Jul 3.09	7 28-Jul 2.26
8 29-Jul 4.96	8 29-Jul 2.22	8 29-Jul 0.87	8 29-Jul 0.35	8 29-Jul 1.52
9 30-Jul 35.38	9 30-Jul 8.77	9 30-Jul 8.55	9 30-Jul 11.46	9 30-Jul 6.60
10 31-Jul 44.82	10 31-Jul 10.70	10 31-Jul 14.02	10 31-Jul 12.85	10 31-Jul 7.25
11 1-Aug 37.61	11 1-Aug 10.38	11 1-Aug 8.61	11 1-Aug 11.81	11 1-Aug 6.81
12 2-Aug 35.98	12 2-Aug 9.69	12 2-Aug 8.24	12 2-Aug 11.25	12 2-Aug 6.80
13 3-Aug 31.77	13 3-Aug 10.85	13 3-Aug 7.75	13 3-Aug 6.89	13 3-Aug 6.29
14 4-Aug 9.02	14 4-Aug 4.63	14 4-Aug 1.65	14 4-Aug 1.43	14 4-Aug 1.31
15 5-Aug 6.57	15 5-Aug 3.12	15 5-Aug 1.33	15 5-Aug 1.15	15 5-Aug 0.97
16 6-Aug 36.85	16 6-Aug 10.99	16 6-Aug 9.89	16 6-Aug 8.42	16 6-Aug 7.55
17 7-Aug 37.66	17 7-Aug 10.98	17 7-Aug 9.21	17 7-Aug 11.10	17 7-Aug 6.37
18 8-Aug 38.30	18 8-Aug 11.66	18 8-Aug 8.87	18 8-Aug 11.08	18 8-Aug 6.70
19 9-Aug 37.54	19 9-Aug 12.75	19 9-Aug 7.87	19 9-Aug 8.50	19 9-Aug 8.42
20 10-Aug 29.24	20 10-Aug 7.94	20 10-Aug 7.46	20 10-Aug 6.93	20 10-Aug 6.91
21 11-Aug 15.54	21 11-Aug 5.33	21 11-Aug 3.60	21 11-Aug 3.72	21 11-Aug 2.89
22 12-Aug 7.86	22 12-Aug 2.40	22 12-Aug 2.60	22 12-Aug 1.69	22 12-Aug 1.18
23 13-Aug 38.33	23 13-Aug 9.41	23 13-Aug 10.86	23 13-Aug 8.94	23 13-Aug 9.12
24 14-Aug 38.70	24 14-Aug 10.57	24 14-Aug 10.56	24 14-Aug 7.89	24 14-Aug 9.69
25 15-Aug 34.76	25 15-Aug 8.54	25 15-Aug 9.44	25 15-Aug 8.17	25 15-Aug 8.61
26 16-Aug 38.15	26 16-Aug 9.57	26 16-Aug 10.85	26 16-Aug 8.05	26 16-Aug 9.68
27 17-Aug 33.36	27 17-Aug 9.96	27 17-Aug 7.44	27 17-Aug 11.01	27 17-Aug 4.95
28 18-Aug 12.60	28 18-Aug 5.26	28 18-Aug 2.30	28 18-Aug 2.20	28 18-Aug 2.85
29 19-Aug 2.55	29 19-Aug 1.06	29 19-Aug 0.47	29 19-Aug 0.55	29 19-Aug 0.47
30 20-Aug 41.21	30 20-Aug 11.29	30 20-Aug 9.26	30 20-Aug 11.10	30 20-Aug 9.56
31 21-Aug 51.59	31 21-Aug 12.01	31 21-Aug 14.60	31 21-Aug 13.16	31 21-Aug 11.83
32 22-Aug 71.98	32 22-Aug 37.61	32 22-Aug 10.47	32 22-Aug 11.23	32 22-Aug 12.67
33 23-Aug 9.44	33 23-Aug 4.55	33 23-Aug 1.22	33 23-Aug 1.64	33 23-Aug 2.03
34 24-Aug 16.19	34 24-Aug 3.82	34 24-Aug 1.40	34 24-Aug 2.63	34 24-Aug 8.34
35 25-Aug 10.47	35 25-Aug 4.01	35 25-Aug 1.88	35 25-Aug 2.95	35 25-Aug 1.62
36 26-Aug 5.26	36 26-Aug 3.11	36 26-Aug 1.47	36 26-Aug 0.67	36 26-Aug 0.01
37 27-Aug 8.90	37 27-Aug 2.96	37 27-Aug 2.69	37 27-Aug 3.25	37 27-Aug 0.00
38 28-Aug 57.00	38 28-Aug 16.81	38 28-Aug 12.01	38 28-Aug 18.50	38 28-Aug 9.67
39 29-Aug 68.39	39 29-Aug 15.06	39 29-Aug 11.35	39 29-Aug 26.86	39 29-Aug 15.12
40 30-Aug 52.74	40 30-Aug 15.42	40 30-Aug 10.54	40 30-Aug 15.36	40 30-Aug 11.43
41 31-Aug 46.49	41 31-Aug 11.39	41 31-Aug 12.19	41 31-Aug 14.54	41 31-Aug 8.37
42 1-Sep 22.58	42 1-Sep 6.63	42 1-Sep 4.75	42 1-Sep 6.07	42 1-Sep 5.13
43 2-Sep 19.41	43 2-Sep 5.33	43 2-Sep 3.58	43 2-Sep 7.60	43 2-Sep 2.89
44 3-Sep 29.36	44 3-Sep 6.74	44 3-Sep 6.91	44 3-Sep 10.27	44 3-Sep 5.44
45 4-Sep 73.71	45 4-Sep 34.99	45 4-Sep 10.37	45 4-Sep 17.67	45 4-Sep 10.68
46 5-Sep 72.06	46 5-Sep 16.37	46 5-Sep 12.42	46 5-Sep 32.29	46 5-Sep 10.99
47 6-Sep 72.78	47 6-Sep 17.14	47 6-Sep 11.13	47 6-Sep 34.85	47 6-Sep 9.66
48 7-Sep 57.73	48 7-Sep 13.20	48 7-Sep 15.08	48 7-Sep 21.06	48 7-Sep 8.39

- The working directory should have the following files

Name

☐ analyze.R  
☐ functions.R  
☒ clusterwide.txt  
☒ node1.txt  
☒ node2.txt  
☒ node3.txt  
☒ node4.txt

## Running the forecast

### R script HOWTO

How to setup:

- \* Setup R on your computer
- \* put the 'analyze.R' and 'functions.R' scripts in the same folder
- \* call Rscript like this:  
    /usr/bin/Rscript/Rscript analyze.R cpu.txt 96 120 0.99 1 200
- \* NOTE: the location of Rscript on your computer may be different if you're not using OSX. On windows it's C:/Program Files/R/R-3.1.1/bin/Rscript.exe
- \* NOTE: if the input file is not in the directory of the R script, you should use the full path like this:  
    /usr/bin/Rscript/Rscript /full/path/to/analyze.R /full/path/to/cpu.txt 96 120 0.99 1 200

INPUT:

#arg 1 is name of input file  
#arg 2 is threshold value  
#arg 3 is smooth window [default 120] - #days to look back  
#arg 4 is quantile [default .99]  
#arg 5 is capacity locked to past max (1 if true 0 if false)  
#arg 6 is #days or #samples to project out [default is data length / 2]

EXAMPLE:

Rscript analyze.R cpu.txt 96 120 0.99 1 200

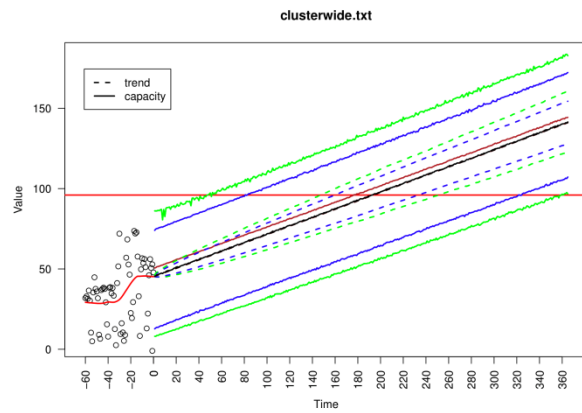
OUTPUT:

PDF of plot (same name as input file except with PDF extension) tab delimited file with HRE at quantile chosen on input [see default above]

LIMITATIONS:

Right now, the program doesn't look at the date column, only the data column and assumes that the dates are increasing one day or sample at a time. So, if this is not the case, it can give some wrong results. This will be fixed in the future.

## Clusterwide headroom expiry at 75% capacity




- 48 days to reach 75% utilization at 99% forecast quantile

```
E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R
clusterwide.txt 96 30 .99 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] 193
$quantiles
      quantile headroom
[1,]      0.99      48
```

- 84 days to reach 75% utilization at 95% forecast quantile

```
E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R
clusterwide.txt 96 30 .95 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] 193
$quantiles
      quantile headroom
[1,]      0.95      84
```

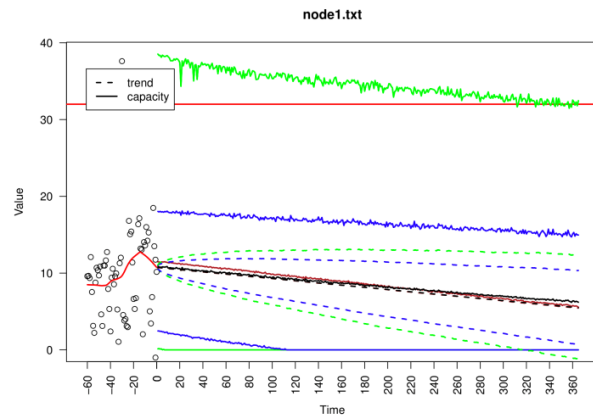
## Clusterwide headroom expiry at 100% capacity



- 165 days to reach 100% utilization at 99% forecast quantile

```
E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R
clusterwide.txt 128 30 .99 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] 314
$quantiles
      quantile headroom
[1,]      0.99      165
```

## Node1 headroom expiry at 100% capacity



- 1 day to reach 100% utilization at 99% forecast quantile

```
E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R node1.txt 32
30 .99 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] NA

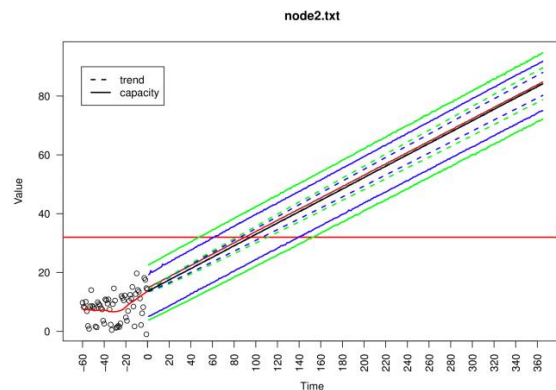
$quantiles
      quantile headroom
[1,]      0.99        1
```

- NA at 95% forecast quantile. This means stable and will not reach the headroom

```
E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R node1.txt 32
30 .95 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] NA

$quantiles
      quantile headroom
[1,]      0.95       NA
```

## Node2 headroom expiry at 100% capacity

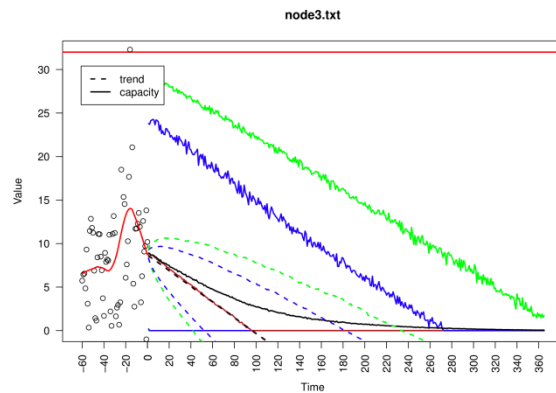


- 48 days to reach 100% utilization at 99% forecast quantile

```
E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R node2.txt 32
30 .99 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] 95

$quantiles
      quantile headroom
[1,]      0.99       48
```

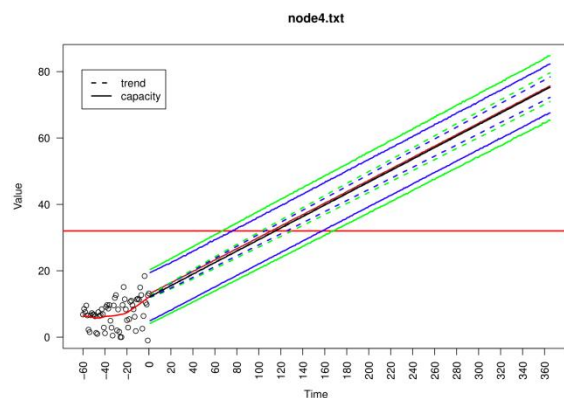
## Node3 headroom expiry at 100% capacity



- NA at 99% forecast quantile. This means stable and will not reach the headroom

```
E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R node3.txt 32
30 .99 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] NA
$quantiles
      quantile headroom
[1,]      0.99      NA
```

## Node4 headroom expiry at 100% capacity



- 67 days to reach 100% utilization at 99% forecast quantile

```
E:\GitHub\forecast examples\monte carlo\tmp>"C:\Program Files\R\R-3.2.2\bin\Rscript.exe" analyze.R node4.txt 32
30 .99 0 365
Running Monte Carlo. Please wait
null device
      1
$mean
[1] 116
$quantiles
      quantile headroom
[1,]      0.99      67
```